

Rahul Sharma

Senior AI & Machine Learning Engineer • rahul.sharma@example.com • +1 (555) 234-5678 • San Francisco, CA

PROFESSIONAL SUMMARY

Innovative Senior AI/ML Engineer with 6.5+ years of experience designing and deploying scalable deep learning pipelines, generative AI solutions, and NLP applications. Proven track record deploying models to AWS using Docker, FastAPI, and Kubernetes.

TECHNICAL SKILLS

Python, Machine Learning, Deep Learning, NLP, LLMs, Generative AI, PyTorch, Scikit-Learn, TensorFlow, Pandas, NumPy, FastAPI, Docker, Kubernetes, AWS, SQL, PostgreSQL, Git, CI/CD, REST APIs, Linux

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer – NeuralScale Inc. (2021 - Present (3+ years))

- Architected NLP and LLM fine-tuning pipelines using PyTorch, Hugging Face, and Scikit-Learn for document comprehension.
- Built asynchronous microservices with FastAPI and Docker, handling 15M+ daily inference queries on AWS.
- Implemented CI/CD pipelines with GitHub Actions for automated model evaluation and container deployments.
- Mentored team of 5 ML engineers in best practices for model tracking, Pandas, and NumPy optimization.

Data Scientist & ML Developer – Cognitive Insights (2018 - 2021 (3 years))

- Engineered predictive machine learning models using Scikit-Learn, XGBoost, and Python.
- Automated ETL pipelines with SQL, PostgreSQL, and Pandas to clean terabyte-scale datasets.
- Collaborated via Git with backend teams to integrate REST APIs for real-time risk scoring.

EDUCATION

M.S. in Computer Science, Stanford University (2018)

B.Tech in Computer Engineering, IIT Delhi (2016)

KEY PROJECTS

Enterprise RAG Assistant: Built retrieval-augmented generation search over 200k documents using PyTorch, LangChain, and FastAPI.

Automated Document Parser: Developed OCR and NLP classification model achieving 96% accuracy.